

1. Eager Loading (Optimizing Queries with `joinedload`)

By default, SQLAlchemy uses **lazy loading**, meaning related objects are not loaded until explicitly accessed. However, in certain cases, you want to optimize by loading everything in one query. This is where **eager loading** comes in with the `joinedload` method.

Example: Eager Load `student` and `Courses` (One-to-Many)

```
python
Copy code
from sqlalchemy.orm import joinedload

# Eager load the courses when querying students
students_with_courses =
session.query(Student).options(joinedload(Student.courses)).all()

for student in students_with_courses:
    print(f"Student: {student.name}")
    for course in student.courses:
        print(f"  Course: {course.course_name}")
```

- **Eager Loading** prevents multiple round-trips to the database by loading all related data at once.

2. Bulk Insert (Efficiently Insert Multiple Records)

When inserting a large number of records, using `session.add_all()` or `session.add()` for each record can be inefficient. SQLAlchemy provides **bulk insert** methods for improved performance.

Example: Bulk Insert Multiple Students

```
python
Copy code
students = [
    {"name": "George", "age": 20, "grade": "A"},
    {"name": "Helen", "age": 22, "grade": "B"},
    {"name": "Isaac", "age": 21, "grade": "C"},
]

# Bulk insert using SQLAlchemy's built-in method
session.bulk_insert_mappings(Student, students)
session.commit()
```

- **Bulk Insert** is useful for inserting a large number of records efficiently.

3. Advanced Querying (Filtering, Joins, and Aggregates)

SQLAlchemy supports complex queries similar to SQL, such as filtering, joins, and aggregate functions.

Example 1: Filtering with Multiple Conditions

```
python
Copy code
# Query students who are older than 21 and have a grade of 'A'
students = session.query(Student).filter(Student.age > 21, Student.grade == 'A').all()

for student in students:
    print(student)
```

Example 2: Joining Tables

```
python
Copy code
# Join Students and Courses
result = session.query(Student.name,
Course.course_name).join(Student.courses).all()

for student_name, course_name in result:
    print(f"{student_name} is enrolled in {course_name}")
```

Example 3: Aggregating with COUNT and GROUP BY

```
python
Copy code
# Count how many students are enrolled in each course
course_counts = session.query(Course.course_name,
func.count(Student.id)).join(Student.courses).group_by(Course.course_name).all()

for course_name, count in course_counts:
    print(f"{course_name}: {count} students")
```

- **Advanced querying** allows you to build complex queries similar to raw SQL but with the ORM's safety and convenience.

4. Transaction Management (Commit, Rollback, and Savepoint)

Transactions are crucial for maintaining data integrity. SQLAlchemy supports transactions via `session.commit()` and `session.rollback()`. For more complex cases, you can also create **savepoints**.

Example: Basic Transaction with Commit and Rollback

```
python
Copy code
try:
    new_student = Student(name='Jason', age=24, grade='A')
    session.add(new_student)
    session.commit() # Commit the transaction
except:
    session.rollback() # Roll back if an error occurs
    print("Transaction rolled back.")
```

Example: Using Savepoints

```
python
Copy code
from sqlalchemy import exc

# Start a new transaction
session.begin()

# Add the first student and create a savepoint
session.add(Student(name='Kevin', age=23, grade='B'))
session.flush() # Flush the transaction to the database
savepoint = session.begin_nested() # Create a savepoint

try:
    # Add another student that could cause an error
    session.add(Student(name='Invalid Student', age=-1, grade='C'))
    session.flush() # Force the error here
except exc.SQLAlchemyError:
    print("Rolling back to savepoint.")
    savepoint.rollback() # Rollback to the savepoint (does not affect
    Kevin's addition)

session.commit() # Commit the transaction after resolving the issue
```

5. Composite Primary Keys (Multi-Column Primary Keys)

In some tables, you may need to define a composite primary key (using more than one column to uniquely identify a row).

Example: Composite Primary Key for Enrollment

```
python
Copy code
class Enrollment(Base):
    __tablename__ = 'enrollments'

    student_id = Column(Integer, ForeignKey('students.id'),
primary_key=True)
    course_id = Column(Integer, ForeignKey('courses.id'), primary_key=True)
    enrollment_date = Column(String)

    student = relationship("Student", back_populates="enrollments")
    course = relationship("Course", back_populates="enrollments")
```

- Here, the **student_id** and **course_id** together form a composite primary key for the Enrollment table.

6. Inheritance (Single Table and Multi-Table)

SQLAlchemy supports **inheritance**, which allows different classes to be mapped to a single table or multiple tables.

Example: Single Table Inheritance

```
python
Copy code
class Person(Base):
    __tablename__ = 'persons'

    id = Column(Integer, primary_key=True)
    name = Column(String)
    type = Column(String)

    __mapper_args__ = {
        'polymorphic_identity': 'person',
        'polymorphic_on': type
    }

class Teacher(Person):
    __mapper_args__ = {
        'polymorphic_identity': 'teacher'
    }

class Student(Person):
    __mapper_args__ = {
        'polymorphic_identity': 'student'
    }
```

- **Single Table Inheritance** stores all data in one table and uses a `type` column to distinguish between different classes.
-

7. Custom SQL (Executing Raw SQL Queries)

Sometimes, you may need to run raw SQL queries directly.

Example: Executing Raw SQL

```
python
Copy code
# Executing raw SQL queries
result = session.execute("SELECT * FROM students WHERE age > :age", {'age': 20})

for row in result:
    print(row)
```

- This is useful for complex queries that are difficult to express with the ORM.

8. Declarative Mixins

SQLAlchemy supports **mixins** to share common logic or fields across models.

Example: Timestamp Mixin

```
python
Copy code
from datetime import datetime

class TimestampMixin:
    created_at = Column(DateTime, default=datetime.utcnow)
    updated_at = Column(DateTime, default=datetime.utcnow,
onupdate=datetime.utcnow)

class Student(TimestampMixin, Base):
    __tablename__ = 'students'

    id = Column(Integer, primary_key=True)
    name = Column(String)
    age = Column(Integer)
    grade = Column(String)
```

- **Mixins** allow you to add common fields like `created_at` and `updated_at` to multiple models.
-

9. Database Events and Triggers

You can define **event listeners** or **triggers** in SQLAlchemy to automatically run code when certain events occur, like inserts or updates.

Example: Listening for Insert Events

```
python
Copy code
from sqlalchemy import event

def before_insert_listener(mapper, connection, target):
    print(f>About to insert {target.name}")

# Attaching the event listener to the Student class
event.listen(Student, 'before_insert', before_insert_listener)
```

- This event will be triggered before a `Student` is inserted.